

## 14-1: Intro to Constraints; NOT NULL and UNIQUE Constraints

1. What is a “constraint” as it relates to data integrity?

Constraints enforce rules that maintain the consistency and reliability of data. A constraint can prevent duplicate values and ensure that a column cannot contain NULL values, or make sure that values in one table correspond to values in another table.

2. What are the limitations of constraints that may be applied at the column level and at the table level?

**Column-level constraints:** These constraints apply to a specific column in the table. They define rules such as NOT NULL, UNIQUE, or CHECK. The limitations include that the constraint can only be applied to a single column at a time.

**Table-level constraints:** These constraints are defined at the table level and can involve multiple columns, such as primary keys, foreign keys, and composite unique constraints. They are more flexible because they can span multiple columns but need to be defined after all the columns are listed in the CREATE TABLE statement.

3. Why is it important to give meaningful names to constraints?

Giving meaningful names to constraints helps with maintenance and error. When errors occur, constraints make it easier to understand which rule was violated.

4. Based on the information provided by the owners, choose a datatype for each column. Indicate the length, precision, and scale for each NUMBER datatype.

Column Name	Data Type	Length/Precision/Scale
id	NUMBER	6
name	VARCHAR2	100
date_opened	DATE	
address	VARCHAR2	200
city	VARCHAR2	100
zip/postal_code	VARCHAR2	10
phone	VARCHAR2	15
email	VARCHAR2	100
manager_id	NUMBER	6
emergency_contact	VARCHAR2	100

5. Use “nullable” to indicate those columns that can have null values.

**Nullable Columns:** emergency\_contact and manager\_id (if no manager is assigned)

**Non-Nullable Columns:** id (Primary key, must be unique and not null), name, date\_opened, address, city, zip/postal\_code, phone, and email (must be unique)

6. Write the CREATE TABLE statement for the Global Fast Foods locations table to define the constraints at the column level.

```
CREATE TABLE global_locations (
    id NUMBER(6) PRIMARY KEY,
    name VARCHAR2(100) NOT NULL,
    date_opened DATE NOT NULL,
    address VARCHAR2(200) NOT NULL,
    city VARCHAR2(100) NOT NULL,
    zip_postal_code VARCHAR2(10),
    phone VARCHAR2(15),
    email VARCHAR2(100) UNIQUE NOT NULL,
    manager_id NUMBER(6),
    emergency_contact VARCHAR2(100),
    CONSTRAINT f_manager FOREIGN KEY (manager_id) REFERENCES
employees(employee_id)
);
```

7. Execute the CREATE TABLE statement in Oracle Application Express.

```
1  CREATE TABLE global_locations (
2      id NUMBER(6) PRIMARY KEY,
3      name VARCHAR2(100) NOT NULL,
4      date_opened DATE NOT NULL,
5      address VARCHAR2(200) NOT NULL,
6      city VARCHAR2(100) NOT NULL,
7      zip_postal_code VARCHAR2(10),
8      phone VARCHAR2(15),
9      email VARCHAR2(100) UNIQUE NOT NULL,
10     manager_id NUMBER(6),
11     emergency_contact VARCHAR2(100),
12     CONSTRAINT f_manager FOREIGN KEY (manager_id) REFERENCES employees(employee_id)
13 );
```

Results Explain Describe Saved SQL History

Table created.

8. Execute a DESCRIBE command to view the Table Summary information.

```

1  DESCRIBE global_locations;
2
3

```

Results Explain Describe Saved SQL History

Object Type TABLE (?) Object GLOBAL\_LOCATIONS (?)

Table	Column	Data Type	Length	Precision	Scale	Primary Key	Nullable	Default	Comment
GLOBAL_LOCATIONS	ID	NUMBER	-	6	0	1	-	-	-
	NAME	VARCHAR2	100	-	-	-	-	-	-
	DATE_OPENED	DATE	7	-	-	-	-	-	-
	ADDRESS	VARCHAR2	200	-	-	-	-	-	-

9. Rewrite the CREATE TABLE statement for the Global Fast Foods locations table to define the UNIQUE constraints at the table level. Do not execute this statement.

```

CREATE TABLE global_locations (
    id NUMBER(6) PRIMARY KEY,
    name VARCHAR2(100) NOT NULL,
    date_opened DATE NOT NULL,
    address VARCHAR2(200) NOT NULL,
    city VARCHAR2(100) NOT NULL,
    zip_postal_code VARCHAR2(10),
    phone VARCHAR2(15),
    email VARCHAR2(100) NOT NULL,
    manager_id NUMBER(6),
    emergency_contact VARCHAR2(100),
    CONSTRAINT email_unique UNIQUE (email),
    CONSTRAINT f_manager FOREIGN KEY (manager_id) REFERENCES
employees(employee_id)
);

```

#### 14-2: PRIMARY KEY, FOREIGN KEY, and CHECK Constraints

1. What is the purpose of a:
  - a. PRIMARY KEY: Uniquely identifies each row in a table.
  - b. FOREIGN KEY: Establishes a relationship between two tables by ensuring that the value in the foreign key column matches a value in the table's primary key or unique key.
  - c. CHECK CONSTRAINT: Used to limit the range of values that can be stored in a column to ensure data validity.

2. Using the column information for the animals table below, name constraints where applicable at the table level, otherwise name them at the column level. Define the primary key (animal\_id). The license\_tag\_number must be unique. The admit\_date and vaccination\_date columns cannot contain null values.

**Column level constraints:**

```
animal_id NUMBER(6) PRIMARY KEY,  
name VARCHAR2(25),  
license_tag_number NUMBER(10) UNIQUE,  
admit_date DATE NOT NULL,  
adoption_id NUMBER(5),  
vaccination_date DATE NOT NULL
```

**Table level constraints:**

```
CREATE TABLE animals (  
    animal_id NUMBER(6) PRIMARY KEY,  
    name VARCHAR2(25),  
    license_tag_number NUMBER(10) UNIQUE,  
    admit_date DATE NOT NULL,  
    adoption_id NUMBER(5),  
    vaccination_date DATE NOT NULL  
);
```

3. Create the animals table. Write the syntax you will use to create the table.

```

1  CREATE TABLE animals (
2      animal_id NUMBER(6) PRIMARY KEY,
3      name VARCHAR2(25),
4      license_tag_number NUMBER(10) UNIQUE,
5      admit_date DATE NOT NULL,
6      adoption_id NUMBER(5),
7      vaccination_date DATE NOT NULL
8  );
9
10
11

```

**Results   Explain   Describe   Saved SQL   History**

Table created.

4. Enter one row into the table. Execute a SELECT \* statement to verify your input. Refer to the graphic below for input.

```

1  INSERT INTO animals (animal_id, name, license_tag_number, admit_date, adoption_id, vaccination_date)
2  VALUES (101, 'Spot', 35540, TO_DATE('2004-10-10', 'YYYY-MM-DD'), 205, TO_DATE('2004-12-12', 'YYYY-MM-DD'));
3
4
5  SELECT * FROM animals WHERE animal_id = 101;
6
7
8

```

**Results   Explain   Describe   Saved SQL   History**

ANIMAL_ID	NAME	LICENSE_TAG_NUMBER	ADMIT_DATE	ADOPTION_ID	VACCINATION_DATE
101	Spot	35540	10-Oct-2004	205	12-Dec-2004

1 rows returned in 0.00 seconds   [Download](#)

5. Write the syntax to create a foreign key (adoption\_id) in the animals table that has a corresponding primary- key reference in the adoptions table. Show both the column-level and

table-level syntax. Note that because you have not actually created an adoptions table, no adoption\_id primary key exists, so the foreign key cannot be added to the animals table.

**Column level:**

```
CREATE TABLE animals (
    animal_id NUMBER(6) PRIMARY KEY,
    name VARCHAR2(25),
    license_tag_number NUMBER(10) UNIQUE,
    admit_date DATE NOT NULL,
    adoption_id NUMBER(5),
    vaccination_date DATE NOT NULL,
    CONSTRAINT fk_adoption FOREIGN KEY (adoption_id) REFERENCES
    adoptions(adoption_id)
);
```

**Table level:**

```
CREATE TABLE animals (
    animal_id NUMBER(6) PRIMARY KEY,
    name VARCHAR2(25),
    license_tag_number NUMBER(10) UNIQUE,
    admit_date DATE NOT NULL,
    adoption_id NUMBER(5),
    vaccination_date DATE NOT NULL,
    CONSTRAINT fk_adoption FOREIGN KEY (adoption_id) REFERENCES
    adoptions(adoption_id)
);
```

6. What is the effect of setting the foreign key in the ANIMAL table as:

- ON DELETE CASCADE: When a row in the parent table (adoptions) is deleted, all corresponding rows in the child table (animals) will also be deleted automatically.
- ON DELETE SET NULL: When a row in the parent table (adoptions) is deleted, the adoption\_id in the child table (animals) will be set to NULL, but the corresponding row in the child table will not be deleted.

7. What are the restrictions on defining a CHECK constraint?

The condition in the CHECK constraint must be a logical expression; it cannot reference other rows or tables, as CHECK constraints only operate on individual rows of the table. A CHECK constraint also cannot contain subqueries and the logic in the CHECK constraint is not enforced on NULL values, so if a column has NULL, it is considered valid.

### 14-3: Managing Constraints

1. What are four functions that an ALTER statement can perform on constraints?
  - ADD
  - DROP
  - DISABLE
  - ENABLE
2. Since the tables are copies of the original tables, the integrity rules are not passed onto the new tables; only the column datatype definitions remain. You will need to add a PRIMARY KEY constraint to the copy\_d\_clients table. Name the primary key copy\_d\_clients\_pk . What is the syntax you used to create the PRIMARY KEY constraint to the copy\_d\_clients.table?

```
1  ALTER TABLE copy_d_clients
2  ADD CONSTRAINT copy_d_clients_pk PRIMARY KEY (client_number);
3
4
```

Results   Explain   Describe   Saved SQL   History

Table altered.

3. Create a FOREIGN KEY constraint in the copy\_d\_events table. Name the foreign key copy\_d\_events\_fk. This key references the copy\_d\_clients table client\_number column. What is the syntax you used to create the FOREIGN KEY constraint in the copy\_d\_events table?

```
1 ALTER TABLE copy_d_events
2 ADD CONSTRAINT copy_d_events_fk FOREIGN KEY (client_number)
3 REFERENCES copy_d_clients(client_number);
4
5
```

Results Explain Describe Saved SQL History

Table altered.

4. Use a SELECT statement to verify the constraint names for each of the tables. Note that the table names must be capitalized.

- a. The constraint name for the primary key in the copy\_d\_clients table is \_\_\_\_\_.

```
1 SELECT constraint_name
2 FROM user_constraints
3 WHERE table_name = 'COPY_D_CLIENTS' AND constraint_type = 'P';
4
5
6
7
8
```

Results Explain Describe Saved SQL History

**CONSTRAINT\_NAME**

COPY\_D\_CLIENTS\_PK

1 rows returned in 3.59 seconds [Download](#)

- b. The constraint name for the foreign key in the copy\_d\_events table is \_\_\_\_\_.

```
1  SELECT constraint_name
2  FROM user_constraints
3  WHERE table_name = 'COPY_D_EVENTS' AND constraint_type = 'R';
4
5
```

Results

Explain

Describe

Saved SQL

History

CONSTRAINT_NAME
COPY_D_EVENTS_FK

1 rows returned in 1.71 seconds      [Download](#)

5. Drop the PRIMARY KEY constraint on the copy\_d\_clients table. Explain your results.

```
1 ALTER TABLE copy_d_clients  
2 DROP CONSTRAINT copy_d_clients_pk;  
3 |  
4  
5
```

Results Explain Describe Saved SQL History

Error at Line 2/17: ORA-02273: this unique/primary key is referenced by some foreign keys  
ORA-06512: at "SYS.WWV\_DBMS\_SQL\_APEX\_220200", line 828  
ORA-06512: at "SYS.DBMS\_SYS\_SQL", line 1658  
ORA-06512: at "SYS.WWV\_DBMS\_SQL\_APEX\_220200", line 813  
ORA-06512: at "APEX\_220200.WWV\_FLOW\_DYNAMIC\_EXEC", line 2046

1. ALTER TABLE copy\_d\_clients  
2. DROP CONSTRAINT copy\_d\_clients\_pk;

0.01 seconds

You cannot drop a primary key if there are foreign keys referencing it, unless those foreign keys are also removed.

6. Add the following event to the copy\_d\_events table. Explain your results.

```
1  INSERT INTO copy_d_events (id, name, event_date, description, cost, venue_id, package_code, theme_code, client_
2    VALUES (140, 'Cline Bas Mitzvah', TO_DATE('15-Jul-2004', 'DD-Mon-YYYY'), 'Church and Private Home', 4500, 105_
3_
4_
5_
6
```

Results Explain Describe Saved SQL History

```
ORA-02291: integrity constraint (US_A296_SQL_S36_ADMIN.COPY_D_EVENTS_FK)
violated - parent key not found
ORA-06512: at "SYS.DBMS_SQL", line 1721
```

```
1. INSERT INTO copy_d_events (id, name, event_date, description, cost,
venue_id, package_code, theme_code, client_number)
2. VALUES (140, 'Cline Bas Mitzvah', TO_DATE('15-Jul-2004', 'DD-MM-YYYY'),
'Church and Private Home', 4500, 105, 87, 77, 7125);
```

You cannot add the event to the table because there is no parent key found.

7. Create an ALTER TABLE query to disable the primary key in the copy\_d\_clients table. Then add the values from #6 to the copy\_d\_events table. Explain your results.

```
1  INSERT INTO copy_d_events (id, name, event_date, description, cost, venue_id, package_code, theme_code, client_
2    VALUES (140, 'Cline Bas Mitzvah', TO_DATE('15-Jul-2004', 'DD-Mon-YYYY'), 'Church and Private Home', 4500, 105_
3_
4
```

Results Explain Describe Saved SQL History

```
1 row(s) inserted.
```

This INSERT statement adds a new row into the copy\_d\_events table.

8. Repeat question 6: Insert the new values in the copy\_d\_events table. Explain your results.

This INSERT statement adds a new row into the copy\_d\_events table.

9. Enable the primary-key constraint in the copy\_d\_clients table. Explain your results.

```
1 ALTER TABLE copy_d_clients ENABLE CONSTRAINT copy_d_clients_pk;
2 |
```

Results Explain Describe Saved SQL History

Table altered.

Enabling the primary key constraint re-applies the uniqueness and NOT NULL to the client\_number column.

10. If you wanted to enable the foreign-key column and reestablish the referential integrity between these two tables, what must be done?

You need to enable the foreign key constraint in the copy\_d\_events table and make sure that all existing rows in the copy\_d\_events table have valid client\_number values that exist in the copy\_d\_clients table.

11. Why might you want to disable and then re-enable a constraint?

When loading large amounts of data, disabling constraints temporarily can speed up the process, and re-enabling them ensures integrity after the data load. Disabling and re-enabling a constraint may also improve performance during bulk operations and ensure data consistency.

12. Query the data dictionary for some of the constraints that you have created. How does the data dictionary identify each constraint type?

The data dictionary identifies each constraint type by labeling it: primary key (P) or check (C).

```
1 SELECT constraint_name, constraint_type
2 FROM user_constraints
3 WHERE table_name = 'COPY_D_CLIENTS';
4 |
```

Results

Explain

Describe

Saved SQL

History

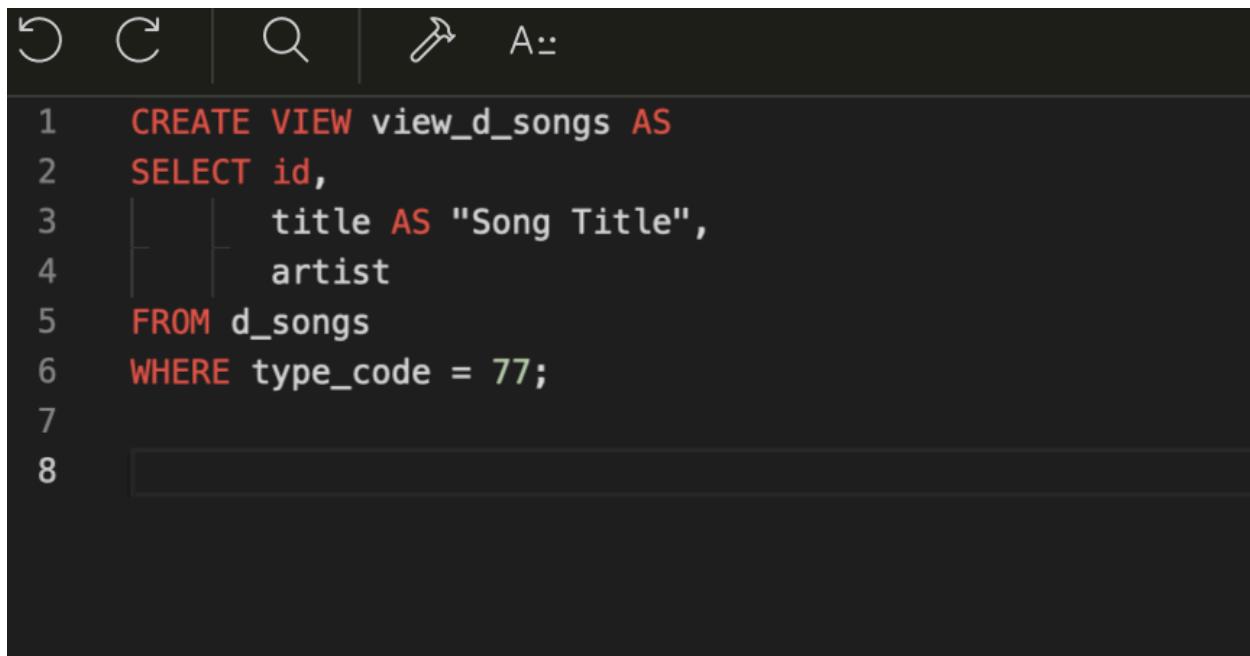
CONSTRAINT_NAME	CONSTRAINT_TYPE
SYS_C0014557705	C
SYS_C0014557706	C
SYS_C0014557707	C
COPY_D_CLIENTS_PK	P

4 rows returned in 1.68 seconds

[Download](#)

## 15-1: Creating Views

1. What are three uses for a view from a DBA's perspective?
  - Security and data restriction.
  - Simplifying complex queries
  - Data aggregation and reporting.
2. Create a simple view called view\_d\_songs that contains the ID, title, and artist from the DJs on Demand table for each "New Age" type code. In the subquery, use the alias "Song Title" for the title column.



The screenshot shows a terminal window with a dark background and light-colored text. At the top, there are several icons: a circular arrow, a 'C' with a horizontal line, a magnifying glass, a wrench, and 'A::'. Below these are two rows of numbers and SQL code. The first row has numbers 1 through 8. The second row contains the SQL code for creating a view:

```
1 CREATE VIEW view_d_songs AS
2 SELECT id,
3     title AS "Song Title",
4     artist
5 FROM d_songs
6 WHERE type_code = 77;
7
8
```

3. 

```
SELECT *
FROM view_d_songs.
```

What was returned?

```
1 SELECT *
2 FROM view_d_songs
3
```

Results Explain Describe Saved SQL History

ID	Song Title	ARTIST
47	Hurrah for Today	The Jubilant Trio
49	Lets Celebrate	The Celebrants

2 rows returned in 0.00 seconds    Download

4. REPLACE view\_d\_songs. Add type\_code to the column list. Use aliases for all columns.

```
1 CREATE OR REPLACE VIEW view_d_songs AS
2 SELECT id AS "Song ID",
3       title AS "Song Title",
4       artist AS "Song Artist",
5       type_code AS "Song Type"
6   FROM d_songs
7 WHERE type_code = 77;
8
9
```

Results Explain Describe Saved SQL History

View created.

5. Jason Tsang, the disk jockey for DJs on Demand, needs a list of the past events and those planned for the coming months so he can make arrangements for each event's equipment setup. As the company manager, you do not want him to have access to the price that clients paid for their events. Create a view for Jason to use that

displays the name of the event, the event date, and the theme description. Use aliases for each column name.

```
1 CREATE VIEW view_jason_events AS
2 SELECT name AS "Event Name",
3       event_date AS "Event Date",
4       description AS "Theme Description"
5 FROM d_events
6 WHERE event_date >= SYSDATE;
7
```

Results

Explain

Describe

Saved SQL

History

View created.

6. It is company policy that only upper-level management be allowed access to individual employee salaries. The department managers, however, need to know the minimum, maximum, and average salaries, grouped by department. Use the Oracle database to prepare a view that displays the needed information for department managers.

```
1 CREATE VIEW view_salary_summary AS
2 SELECT department_id,
3       MIN(salary) AS "Min Salary",
4       MAX(salary) AS "Max Salary",
5       AVG(salary) AS "Avg Salary"
6 FROM employees
7 GROUP BY department_id;
8
9
```

Results

Explain

Describe

Saved SQL

History

View created.

## 15-2: DML Operations and Views

1. Query the data dictionary USER\_UPDATABLE\_COLUMNS to make sure the columns in the base tables will allow UPDATE, INSERT, or DELETE. Use a SELECT statement. All table names in the data dictionary are stored in uppercase.

The screenshot shows a database interface with a SQL editor at the top containing the following code:

```
1  SELECT table_name, column_name, updatable
2  FROM user_updatable_columns
3  WHERE table_name IN ('COPY_D_SONGS', 'COPY_D_EVENTS', 'COPY_D_CDS', 'COPY_D_CLIENTS');
4
5
6
```

Below the editor is a results grid with the following data:

Results	Explain	Describe	Saved SQL	History
TABLE_NAME	COLUMN_NAME	UPDATABLE		
COPY_D_CDS	CD_NUMBER	YES		
COPY_D_CDS	TITLE	YES		
COPY_D_CDS	PRODUCER	YES		
COPY_D_CDS	YEAR	YES		
COPY_D_CLIENTS	CLIENT_NUMBER	YES		
COPY_D_CLIENTS	FIRST_NAME	YES		
COPY_D_CLIENTS	LAST_NAME	YES		
COPY_D_CLIENTS	PHONE	YES		
COPY_D_CLIENTS	EMAIL	YES		

2. Use the CREATE or REPLACE option to create a view of all the columns in the copy\_d\_songs table called view\_copy\_d\_songs.

```
CREATE OR REPLACE VIEW view_copy_d_songs AS
SELECT * FROM copy_d_songs;
```

View created.

3. Use view\_copy\_d\_songs to INSERT the following data into the underlying copy\_d\_songs table. Execute a SELECT \* from copy\_d\_songs to verify your DML command.

```
INSERT INTO view_copy_d_songs (id, title, duration, artist, type_code)
VALUES (88, 'Mello Jello', 2, 'The What', 4);
```

1 row(s) inserted.

The screenshot shows a database interface with a dark theme. At the top, there is a toolbar with icons for back, forward, search, and other functions. Below the toolbar, a code editor window displays the following SQL query:

```
1  SELECT * FROM copy_d_songs;
```

Below the code editor is a results table. The table has a header row with columns: ID, TITLE, DURATION, ARTIST, and TYPE\_CODE. The data rows are:

ID	TITLE	DURATION	ARTIST	TYPE_CODE
45	Its Finally Over	5 min	The Hobbits	12
46	Im Going to Miss My Teacher	2 min	Jane Pop	12
47	Hurrah for Today	3 min	The Jubilant Trio	77
48	Meet Me At the Altar	6 min	Bobby West	1
49	Lets Celebrate	8 min	The Celebrants	77
50	All These Years	10 min	Diana Crooner	88
88	Mello Jello	2	The What	4

At the bottom of the results table, it says "7 rows returned in 0.01 seconds" and has a "Download" link.

4. Create a view based on the DJs on Demand COPY\_D\_CDS table. Name the view read\_copy\_d\_cds. Select all columns to be included in the view. Add a WHERE clause to restrict the year to 2000. Add the WITH READ ONLY option.

The screenshot shows a database interface with a dark theme. At the top, there are several icons: a refresh circle, a search magnifying glass, a refresh arrow, and a font size A±. Below the toolbar, a SQL code editor contains the following script:

```
1 CREATE OR REPLACE VIEW read_copy_d_cds AS
2 SELECT *
3 FROM copy_d_cds
4 WHERE year = 2000
5 WITH READ ONLY;
6
```

Below the code editor, there is a navigation bar with tabs: **Results**, Explain, Describe, Saved SQL, and History. The **Results** tab is currently selected. Underneath the tabs, the message "View created." is displayed.

5. Using the read\_copy\_d\_cds view, execute a DELETE FROM read\_copy\_d\_cds WHERE cd\_number = 90;

The screenshot shows a SQL developer interface. In the top-left corner, there are icons for refresh, search, and edit. Below them, a code editor window contains the following SQL command:

```
1. DELETE FROM read_copy_d_cds WHERE cd_number = 90;
```

Below the code editor, a navigation bar has tabs for 'Results' (which is selected), 'Explain', 'Describe', 'Saved SQL', and 'History'. A yellow callout box highlights the error message and the failed query:

Error at line 1/13: ORA-42399: cannot perform a DML operation on a read-only view  
ORA-06512: at "SYS.WWV\_DBMS\_SQL\_APEX\_220200", line 828  
ORA-06512: at "SYS.DBMS\_SYS\_SQL", line 1658  
ORA-06512: at "SYS.WWV\_DBMS\_SQL\_APEX\_220200", line 813  
ORA-06512: at "APEX\_220200.WWV\_FLOW\_DYNAMIC\_EXEC", line 2046

1. DELETE FROM read\_copy\_d\_cds WHERE cd\_number = 90;

6. Use REPLACE to modify read\_copy\_d\_cds. Replace the READ ONLY option with WITH CHECK OPTION CONSTRAINT ck\_read\_copy\_d\_cds. Execute a SELECT \* statement to verify that the view exists.

```
1 CREATE OR REPLACE VIEW read_copy_d_cds AS
2 SELECT *
3 FROM copy_d_cds
4 WHERE year = 2000
5 WITH CHECK OPTION CONSTRAINT ck_read_copy_d_cds;
6
7
```

**Results**   [Explain](#)   [Describe](#)   [Saved SQL](#)   [History](#)

View created.

0.02 seconds

Language SQL Rows 10 Clear Command Find Tables Save Run

1 SELECT \*  
2 FROM read\_copy\_d\_cds;  
3  
4

Results Explain Describe Saved SQL History

CD_NUMBER	TITLE	PRODUCER	YEAR
91	Party Music for All Occasions	The Music Man	2000
94	Carpe Diem	R & B Inc.	2000

2 rows returned in 0.00 seconds [Download](#)

7. Use the read\_copy\_d\_cds view to delete any CD of year 2000 from the underlying copy\_d\_cds.

1 DELETE FROM read\_copy\_d\_cds  
2 WHERE year = 2000;  
3  
4

Results Explain Describe Saved SQL History

2 row(s) deleted.

0.00 seconds

8. Use the read\_copy\_d\_cds view to delete cd\_number 90 from the underlying copy\_d\_cds table.

```
↻ C | Q | ⚒ A::  
1 DELETE FROM read_copy_d_cds  
2 WHERE cd_number = 90;  
3 |  
4  
  
Results Explain Describe Saved SQL History  
  
0 row(s) deleted.
```

9. Use the read\_copy\_d\_cds view to delete year 2001 records.

```
↻ C | Q | ⚒ A::  
1 DELETE FROM read_copy_d_cds  
2 WHERE year = 2001;  
3 |  
4  
5  
  
Results Explain Describe Saved SQL History  
  
0 row(s) deleted.  
  
0.00 seconds
```

10. Execute a SELECT \* statement for the base table copy\_d\_cds. What rows were deleted?

The screenshot shows a database interface with a code editor at the top containing the SQL query:

```
1 SELECT * FROM copy_d_cds;
```

Below the code editor is a results table with the following data:

CD_NUMBER	TITLE	PRODUCER	YEAR
90	The Celebrants Live in Concert	Old Town Records	1997
92	Back to the Shire	Middle Earth Records	2002
93	Songs from My Childhood	Old Town Records	1999
95	Here Comes the Bride	The Music Man	2001
96	Graduation Songbook	Tunes Are Us	1998
98	Whirled Peas	Old Town Records	2004

At the bottom of the results table, it says "6 rows returned in 0.01 seconds" and has a "Download" link.

11. What are the restrictions on modifying data through a view?

Not all views are updatable. For a view to be updatable, it must meet certain conditions. WITH READ ONLY it prevents any modifications to the underlying table through the view and WITH CHECK OPTION it makes sure that only rows that meet the view's WHERE clause can be inserted or updated. Any operation that violates the view's condition is rejected.

12. What is Moore's Law? Do you consider that it will continue to apply indefinitely? Support your opinion with research from the internet.

Moore's Law is the idea that the number of transistors on a computer chip doubles approximately every two years, which leads to faster and cheaper computers over time. However, as transistors get smaller and approach the size of atoms, it becomes harder to keep doubling their number. This means that Moore's Law might not continue indefinitely because of physical limits. To keep improving technology, scientists and engineers are exploring new ideas, like 3D chips or quantum computing, which could help maintain progress even after Moore's

Law slows down. In short, while Moore's Law has been a reliable guide for technological growth, it may not hold forever, and future advances may rely on new technologies.

<https://cap.csail.mit.edu/death-moores-law-what-it-means-and-what-might-fill-gap-going-forward>

13. What is the “singularity” in terms of computing?

In terms of computing, the singularity refers to the hypothetical future point where technological growth becomes uncontrollable and irreversible, leading to unforeseeable changes in human civilization.

## 15-3: Managing Views

1. Create a view from the copy\_d\_songs table called view\_copy\_d\_songs that includes only the title and artist. Execute a SELECT \* statement to verify that the view exists.

The screenshot shows a database interface with a code editor at the top and a results table below. The code editor contains the following SQL:

```
1 CREATE VIEW view_copy_d_songs AS
2 SELECT title, artist
3 FROM copy_d_songs;
4
5 SELECT * FROM view_copy_d_songs;
```

The results table has columns: Results, Explain, Describe, Saved SQL, and History. The Results tab is selected, showing the following data:

ID	TITLE	DURATION	ARTIST	
45	Its Finally Over	5 min	The Hobbits	12
46	Im Going to Miss My Teacher	2 min	Jane Pop	12
47	Hurrah for Today	3 min	The Jubilant Trio	77
48	Meet Me At the Altar	6 min	Bobby West	1
49	Lets Celebrate	8 min	The Celebrants	77
50	All These Years	10 min	Diana Crooner	88
88	Mello Jello	2	The What	4

At the bottom, it says "7 rows returned in 0.00 seconds" and has a "Download" link.

2. Issue a DROP view\_copy\_d\_songs. Execute a SELECT \* statement to verify that the view has been deleted.

```
↻ ⌂ | ⌂ | ⌂ | A..  
1  DROP VIEW view_copy_d_songs;  
2  
3  SELECT * FROM view_copy_d_songs;  
4  
5  
6  
7
```

Results Explain Describe Saved SQL History

 Error at line 1/15: ORA-00942: table or view does not exist

3. Create a query that selects the last name and salary from the Oracle database. Rank the salaries from highest to lowest for the top three employees.

```
↻ C | Q | ↗ A: 1 SELECT last_name, salary  
2 FROM employees  
3 ORDER BY salary DESC  
4 FETCH FIRST 3 ROWS ONLY;  
5  
6
```

Results   Explain   Describe   Saved SQL   History

LAST_NAME	S
King	24000
Kochhar	17000
De Haan	17000

3 rows returned in 0.00 seconds   [Download](#)

4. Construct an inline view from the Oracle database that lists the last name, salary, department ID, and maximum salary for each department. Hint: One query will need to calculate maximum salary by department ID.

```
1  SELECT last_name, salary, department_id,
2    |      (SELECT MAX(salary) FROM employees WHERE department_id = e.department_id) AS max_salary
3  FROM employees e;
4
5
```

Results Explain Describe Saved SQL History

LAST_NAME	SALARY	DEPARTMENT_ID	MAX_SALARY
Mourgos	5800	50	5800
Rajs	3500	50	5800
Davies	3100	50	5800
Matos	2600	50	5800
Vargas	2500	50	5800
Bell	3500	50	5800
Heiden	2600	50	5800
Higgins	12000	110	12000

5. Create a query that will return the staff members of Global Fast Foods ranked by salary from lowest to highest.

```
1  SELECT * FROM f_staffs
2  ORDER BY salary ASC;
3
4
```

Results   Explain   Describe   Saved SQL   History

ID	FIRST_NAME	LAST_NAME	BIRTHDATE	SALARY	OVERTIME_RATE	TRAINING	STAFF_TYPE	MANAGE
12	Sue	Doe	01-Jul-1980	6.75	10.25	-	Order Taker	19
9	Bob	Miller	19-Mar-1979	10	-	Grill	Cook	19
19	Monique	Tuttle	30-Mar-1969	60	-	-	Manager	-

3 rows returned in 0.01 seconds   [Download](#)

Extension Exercises:

```
1  CREATE TABLE my_departments AS
2  SELECT * FROM departments;
3
4  SELECT * FROM my_departments;
5
```

Results Explain Describe Saved SQL History

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
50	Shipping	124	1500
60	IT	103	1400
80	Sales - Europe	149	2500
85	Sales - Americas	149	2100
90	Executive	100	1700

1.

```
1 DESCRIBE my_departments;
2
3
4
```

Results	Explain	Describe	Saved SQL	History						
Object Type		TABLE	Object		MY_DEPARTMENTS					
Table	Column	Data Type	Length	Precision	Scale	Primary Key	Nullable	Default	Comments	
MY_DEPARTMENTS	DEPARTMENT_ID	NUMBER	-	4	0	-	✓	-	-	
	DEPARTMENT_NAME	VARCHAR2	30	-	-	-	-	-	-	
	MANAGER_ID	NUMBER	-	6	0	-	✓	-	-	
	LOCATION_ID	NUMBER	-	4	0	-	✓	-	-	

2.

```
1  CREATE VIEW view_my_departments AS
2  SELECT department_id, department_name
3  FROM my_departments;
4
5
6
7
```

**Results**   [Explain](#)   [Describe](#)   [Saved SQL](#)   [History](#)

View created.

3.

```
↻ C | Q | ↗ A::  
1 INSERT INTO view_my_departments (department_id, department_name)  
2 VALUES (105, 'Advertising');  
3  
4 INSERT INTO view_my_departments (department_id, department_name)  
5 VALUES (120, 'Custodial');  
6  
7 INSERT INTO view_my_departments (department_id, department_name)  
8 VALUES (130, 'Planning');  
9  
  
Results Explain Describe Saved SQL History  
  
1 row(s) inserted.  
0.01 seconds
```

4.

```
1  ALTER TABLE my_departments  
2  ADD CONSTRAINT pk_department_id PRIMARY KEY (department_id);  
3
```

Results   Explain   Describe   Saved SQL   History

Table altered.

0.02 seconds

- 5.
6. INSERT INTO view\_my\_departments (department\_name)  
VALUES ('Human Resources');

```
1  INSERT INTO view_my_departments (department_id, department_name)
2  VALUES (220, 'Human Resources');
3
```

Results   Explain   Describe   Saved SQL   History

7. 1 row(s) inserted.

```
1  SELECT * FROM view_my_departments;
2  |
```

Results

Explain

Describe

Saved SQL

History

8.

DEPARTMENT_ID	DEPARTMENT_NAME
10	Administration
20	Marketing
50	Shipping
60	IT
80	Sales - Europe
85	Sales - Americas
90	Executive
110	Accounting

```
1  CREATE OR REPLACE VIEW view_my_departments AS
2  SELECT department_id, department_name, location_id
3  FROM my_departments;
4
```

Results Explain Describe Saved SQL History

View created.

9.

10. ALTER TABLE my\_departments  
MODIFY location\_id NOT NULL;

11. CREATE VIEW view\_location\_departments AS  
SELECT d.department\_name, l.street\_address, l.city, l.state  
FROM locations l  
JOIN departments d ON l.location\_id = d.location\_id  
WHERE l.city IN ('U.S. City Names');

## 16-1: Working with Sequences

1. Using CREATE TABLE AS subquery syntax, create a seq\_d\_songs table of all the columns in the DJs on Demand database table d\_songs. Use the SELECT \* in the subquery to make sure that you have copied all of the columns.

The screenshot shows a SQL interface with a code editor at the top and a results table below. The code editor contains the following SQL:

```
1 CREATE TABLE seq_d_songs AS
2 SELECT * FROM seq_d_songs;
3
4
```

The results table has the following structure and data:

Results	Explain	Describe	Saved SQL	History
ID	TITLE	DURATION	ARTIST	TYPE_CODE
45	Its Finally Over	5 min	The Hobbits	12
46	Im Going to Miss My Teacher	2 min	Jane Pop	12
47	Hurrah for Today	3 min	The Jubilant Trio	77
48	Meet Me At the Altar	6 min	Bobby West	1
49	Lets Celebrate	8 min	The Celebrants	77
50	All These Years	10 min	Diana Crooner	88

At the bottom, it says "6 rows returned in 0.00 seconds" and has a "Download" link.

2. Because you are using copies of the original tables, the only constraints that were carried over were the NOT NULL constraints. Create a sequence to be used with the primary-key column of the seq\_d\_songs table. To avoid assigning primary-key numbers to these tables that already exist, the sequence should start at 100 and have a maximum

value of 1000. Have your sequence increment by 2 and have NOCACHE and NOCYCLE. Name the sequence seq\_d\_songs\_seq.

```
1 CREATE SEQUENCE seq_d_songs_seq
2 START WITH 100
3 INCREMENT BY 2
4 MAXVALUE 1000
5 NOCACHE
6 NOCYCLE;
7 |
8
```

Results Explain Describe Saved SQL History

Sequence created.

3. Query the USERSEQUENCES data dictionary to verify the seq\_d\_songs\_seq SEQUENCE settings.

```
1  SELECT sequence_name  
2  FROM user_sequences  
3  WHERE sequence_name = 'SEQ_D_SONGS_SEQ';  
4  
5
```

Results   Explain   Describe   Saved SQL   History

### SEQUENCE\_NAME

SEQ\_D\_SONGS\_SEQ

1 rows returned in 0.01 seconds      Download

4. Insert two rows into the seq\_d\_songs table. Be sure to use the sequence that you created for the ID column. Add the two songs shown in the graphic.

```
1  INSERT INTO seq_d_songs (id, title, duration, artist, type_code)  
2  VALUES (seq_d_songs_seq.NEXTVAL, 'Island Fever', '5 min', 'Hawaiian Islanders', 12);  
3  
4  INSERT INTO seq_d_songs (id, title, duration, artist, type_code)  
5  VALUES (seq_d_songs_seq.NEXTVAL, 'Castle of Dreams', '4 min', 'The Wanderers', 77);  
6  
7  |
```

Results   Explain   Describe   Saved SQL   History

1 row(s) inserted.

5. Write out the syntax for seq\_d\_songs\_seq to view the current value for the sequence. Use the DUAL table. (Oracle Application Developer will not run this query.)

```
SELECT seq_d_songs_seq.CURRVAL FROM dual;
```

6. What are three benefits of using SEQUENCEs?

Sequences provide unique numbers, ensuring that primary key values do not repeat.

Sequences can also have different values to meet various application needs and are generated fairly quickly.

7. What are the advantages of caching sequence values?

Caching sequence values can reduce the number of operations needed to retrieve sequence values.

8. Name three reasons why gaps may occur in a sequence?

Some reasons why gaps may occur in a sequence is If a roll back applies, the number generated during that transaction can be lost. Another reason is if a sequence is cached and the system crashes, those values can no longer be used. Lastly, when NEXTVAL is called for generating sequence values and the generated value is not used, it can cause gaps.

## Extension Exercises:

```
1  CREATE TABLE students_2 (
2      student_id NUMBER(8) PRIMARY KEY,
3      first_name VARCHAR2(50),
4      last_name VARCHAR2(50),
5      enrollment_date DATE
6  );
7
```

Results Explain Describe Saved SQL History

able created.

1.

```
1 CREATE SEQUENCE student_id_seq
2 START WITH 1
3 INCREMENT BY 1
4 MAXVALUE 99999999
5 NOCACHE
6 NOCYCLE;
7
```

Results

Explain

Describe

Saved SQL

History

Sequence created.

2.

```

1  INSERT INTO students_2 (student_id, first_name, last_name, enrollment_date)
2  VALUES (student_id_seq.NEXTVAL, 'John', 'Doe', SYSDATE);
3
4  INSERT INTO students_2 (student_id, first_name, last_name, enrollment_date)
5  VALUES (student_id_seq.NEXTVAL, 'Jane', 'Smith', SYSDATE);
6
7  SELECT * FROM students_2

```

Results   Explain   Describe   Saved SQL   History

STUDENT_ID	FIRST_NAME	LAST_NAME	ENROLLMENT_DATE
2	Jane	Smith	02-Dec-2024
1	John	Doe	02-Dec-2024

2 rows returned in 0.01 seconds   [Download](#)

3.

## 16-2: Indexes and Synonyms

1. What is an index and what is it used for?

An index is a database object that improves the speed of data retrieval operations on a table. Indexes are mainly used to speed up query performance, especially for columns frequently involved in WHERE, ORDER BY, JOIN, and other operations.

2. What is a ROWID, and how is it used?

A ROWID is a unique identifier for each row in a table. It represents the physical location of a row in a database.

3. When will an index be created automatically?

When a primary key or unique key constraint is defined on a table and when a unique constraint or primary key is added to a column.

4. Create a nonunique index (foreign key) for the DJs on Demand column (cd\_number) in the D\_TRACK\_LISTINGS table. Use the Oracle Application Developer SQL Workshop Data Browser to confirm that the index was created.

```
1 CREATE INDEX idx_cd_number  
2 ON d_track_listings (cd_number);  
3 |  
4
```

Results   Explain   Describe   Saved SQL   History

Index created.

5. Use the join statement to display the indexes and uniqueness that exist in the data dictionary for the DJs on Demand D\_SONGS table.

```
1  SELECT ui.index_name, ui.uniqueness
2  FROM user_indexes ui
3  JOIN user_tab_columns utc ON ui.table_name = utc.table_name
4  WHERE ui.table_name = 'D_SONGS';
5
6
```

Results   Explain   Describe   Saved SQL   History

INDEX_NAME	UNIQUENESS
D_SNG_ID_PK	UNIQUE

5 rows returned in 1.47 seconds   [Download](#)

6. Use a SELECT statement to display the index\_name, table\_name, and uniqueness from the data dictionary USER\_INDEXES for the DJs on Demand D\_EVENTS table.

```
1  SELECT index_name, table_name, uniqueness
2  FROM user_indexes
3  WHERE table_name = 'D_EVENTS';
4
```

Results Explain Describe Saved SQL History

INDEX_NAME	TABLE_NAME	UNIQUENESS
D_EVE_ID_PK	D_EVENTS	UNIQUE

rows returned in 0.04 seconds [Download](#)

7. Write a query to create a synonym called dj\_tracks for the DJs on Demand d\_track\_listings table.

```
1  CREATE SYNONYM dj_tracks FOR d_track_listings;
2
```

Results Explain Describe Saved SQL History

Synonym created.

8. Create a function-based index for the last\_name column in DJs on Demand D\_PARTNERS table that makes it possible not to have to capitalize the table name for searches. Write a SELECT statement that would use this index.

```

1  CREATE INDEX idx_last_name_lower
2  ON d_partners (LOWER(last_name));
3
4  SELECT * FROM d_partners
5  WHERE LOWER(last_name) = 'plumb';
6

```

Results Explain Describe Saved SQL History

ID	FIRST_NAME	LAST_NAME	EXPERTISE	SPECIALTY	AUTH_EXPENSE_AMT	MANAGER_ID	PARTNER_TYPE
33	Allison	Plumb	Event Planning	-	300000	33	Manager

1 rows returned in 0.01 seconds [Download](#)

9. Create a synonym for the D\_TRACK\_LISTINGS table. Confirm that it has been created by querying the data dictionary.

```

1  CREATE SYNONYM track_listings FOR d_track_listings;
2
3  |SELECT * FROM user_synonyms
4  WHERE synonym_name = 'TRACK_LISTINGS';
5
6

```

Results Explain Describe Saved SQL History

SYNONYM_NAME	TABLE_OWNER	TABLE_NAME	DB_LINK	ORIGIN_CON_
TRACK_LISTINGS	US_A296_SQL_S36_ADMIN	D_TRACK_LISTINGS	-	3

1 rows returned in 0.05 seconds [Download](#)

10. Drop the synonym that you created in question 9.

```
1  DROP SYNONYM track_listings;  
2  |
```

**Results**   [Explain](#)   [Describe](#)   [Saved SQL](#)   [History](#)

Synonym dropped.

0.01 seconds

## 17-1: Controlling User Access

1. What are system privileges concerned with?

System privileges enable users to manage the database environment itself, including creating, altering, or dropping objects like tables, views, or users, and managing other database-level operations.

2. What are object privileges concerned with?

Object privileges are concerned with controlling access to specific database objects like tables, views, or procedures.

3. What is another name for object security?

Object-level security.

4. What commands are necessary to allow Scott access to the database with a password of tiger?

CREATE USER scott IDENTIFIED BY tiger;

GRANT CONNECT TO scott;

5. What are the commands to allow Scott to SELECT from and UPDATE the d\_clients table?

GRANT SELECT, UPDATE ON d\_clients TO scott;

6. What is the command to allow everybody the ability to view the d\_songs table?

GRANT SELECT ON d\_songs TO PUBLIC;

7. Query the data dictionary to view the object privileges granted to you the user.

```

1  SELECT * FROM user_tab_privs;
2

```

Results						
GRANTEE	OWNER	TABLE_NAME	GRANTOR	PRIVILEGE	GRANTABLE	HIERARCHY
PUBLIC	SYS	US_A296_SQL_S36_ADMIN	US_A296_SQL_S36_ADMIN	INHERIT PRIVILEGES	NO	NO

1 rows returned in 3.64 seconds    [Download](#)

8. What privilege should a user be given to create tables?

GRANT CREATE TABLE TO username;

9. If you create a table, how can you pass along privileges to other users just to view your table?

GRANT SELECT ON your\_table TO other\_user;

10. What syntax would you use to grant another user access to your copy\_employees table?  
 GRANT SELECT, INSERT, UPDATE ON copy\_employees TO another\_user;

11. How can you find out what privileges you have been granted for columns in the tables belonging to others?

SELECT \* FROM all\_tab\_privs

WHERE grantee = 'your\_username';

## 17-2: Creating and Revoking Object Privileges

1. What is a role?

A named collection of privileges that can be granted to users or other roles.

2. What are the advantages of a role to a DBA?

Roles make it easier to manage permissions, saves time, reduces errors, flexibility in managing permissions across different levels, and ensures that users only get necessary permissions.

3. Give the ability to another user in your class to look at one of your tables. Give him the right to let other students have that ability.

```
GRANT SELECT ON my_table TO student2 WITH GRANT OPTION;
```

4. You are the DBA. You are creating many users who require the same system privileges. What should you use to make your job easier?

I would create roles and assign the required system privileges to these roles. Then, I'd assign the roles to users. This would allow me to manage privileges efficiently.

```
CREATE ROLE role_name;  
GRANT system_privileges TO role_name;  
GRANT role_name TO user_name;
```

5. What is the syntax to accomplish the following?

- Create a role of manager that has the privileges to select, insert, and update and delete from the employees table

```
CREATE ROLE manager;
```

```
GRANT SELECT, INSERT, UPDATE, DELETE ON employees TO manager;
```

- Create a role of clerk that just has the privileges of select and insert on the employees table

```
CREATE ROLE clerk;
```

```
GRANT SELECT, INSERT ON employees TO clerk;
```

- Grant the manager role to user scott

```
GRANT manager TO scott;
```

- Revoke the ability to delete from the employees table from the manager role

```
REVOKE DELETE ON employees FROM manager;
```

6. What is the purpose of a database link?

A database link is used to connect and allow communication between two separate databases.

### 17-3: Regular Expressions

1. Working with the employees table, and using regular expressions, write a query that returns employees whose first names start with a “S” (uppercase) followed by either a “t” (lowercase) or “h” (lowercase).

```
1  SELECT *
2  FROM employees
3  WHERE REGEXP_LIKE(first_name, '^S[t|h]');
4
```

The screenshot shows a database query results interface. At the top, there are tabs for Results, Explain, Describe, Saved SQL, and History. The Results tab is selected. Below the tabs is a SQL query window containing the provided code. The main area displays the results of the query, which are two rows from the employees table. The columns shown are EMPLOYEE\_ID, FIRST\_NAME, LAST\_NAME, EMAIL, PHONE\_NUMBER, HIRE\_DATE, JOB\_ID, SALARY, and COMMISSION\_PCT. The data is as follows:

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY	COMMISSION_PCT
100	Steven	King	SKING	515.123.4567	17-Jun-2002	AD_PRES	24000	-
205	Shelley	Higgins	SHIGGINS	515.123.8080	07-Jun-2009	AC_MGR	12000	-

2 rows returned in 0.01 seconds [Download](#)

2. Investigate the LOCATIONS table.

- a. Describe the table.

```
1 DESCRIBE locations;
2
```

Object Type TABLE ⓘ										Object LOCATIONS ⓘ
Table	Column	Data Type	Length	Precision	Scale	Primary Key	Nullable	Default	Comment	
LOCATIONS	LOCATION_ID	NUMBER	-	4	0	1	-	-	-	
	STREET_ADDRESS	VARCHAR2	40	-	-	-	✓	-	-	
	POSTAL_CODE	VARCHAR2	12	-	-	-	✓	-	-	
	CITY	VARCHAR2	30	-	-	-	-	-	-	
	STATE_PROVINCE	VARCHAR2	25	-	-	-	✓	-	-	
	COUNTRY_ID	NUMBER	-	3	0	-	-	-	-	

- b. Perform a select that returns all rows and all columns of that table.

```
1  SELECT * FROM locations;
2  |
3
```

Results   Explain   Describe   Saved SQL   History

LOCATION_ID	STREET_ADDRESS	POSTAL_CODE	CITY	STATE_PROVINCE	COUNTRY_ID
1800	460 Bloor St. W.	ON M5S 1X8	Toronto	Ontario	2
2500	Magdalen Centre, The Oxford Science Park	OX9 9ZB	Oxford	Oxford	44
1400	2014 Jabberwocky Rd	26192	Southlake	Texas	1
1500	2011 Interiors Blvd	99236	South San Francisco	California	1
1700	2004 Charade Rd	98199	Seattle	Washington	1
2100	Av. Rio Branco	20090-003	Rio de Janeiro	Rio de Janeiro	55

6 rows returned in 0.01 seconds   [Download](#)

- c. Write a query using regular expressions that removes the spaces in the street\_address column in the LOCATIONS table.

```
1  SELECT REGEXP_REPLACE(street_address, ' ', '') AS street_address_no_spaces
2  FROM locations;
3
4
```

Results

Explain

Describe

Saved SQL

History

#### STREET\_ADDRESS\_NO\_SPACES

460BloorSt.W.

MagdalenCentre,TheOxfordSciencePark

2014JabberwockyRd

2011InteriorBlvd

2004CharadeRd

Av.RioBranco

6 rows returned in 0.00 seconds

[Download](#)